

# SUMIT YADAV

Mechanistic Interpretability | AI Alignment & Safety | NLP

076bct088.sumit@pcampus.edu.np | +977-9819856148 | [sumityadav.com.np](https://sumityadav.com.np) | [Google Scholar](https://scholar.google.com/citations?user=076bct088) | [github.com/rockerritesh](https://github.com/rockerritesh) | [LinkedIn](https://www.linkedin.com/in/sumit-yadav)

## RESEARCH STATEMENT

I work at the intersection of **mechanistic interpretability**, **AI alignment**, and **robustness**. My research studies how safety-aligned LLMs fail silently, through over-refusals, geometric misrepresentation, and surface-level triggers, and how the internal structure of representations can be understood and steered to make models safer and more reliable. Much of my strongest work is sole-authored: I combine representation-space analysis, adversarial probing, and systems engineering into empirically grounded methods. My work has been accepted at **ACL 2026 (CORE A\*)**, and I am especially interested in extending interpretability and safety to multilingual and low-resource settings.

## SELECTED RESEARCH & PUBLICATIONS

**SafeConstellations: Mitigating Over-Refusals in LLMs Through Task-Aware Representation Steering** 2026  
ACL 2026 Main (CORE A\*)

- LLMs trace consistent, task-specific “*constellation*” trajectories in embedding space that shift predictably between refusal and non-refusal states.
- **SafeConstellations** is an inference-time, task-aware trajectory-shifting method that nudges representations toward non-refusal pathways *only* on over-refusal-prone tasks, cutting over-refusals by up to **73%** with minimal utility loss and no retraining. (*with U. Maskey, M. Dras, U. Naseem.*)

**Your Intervention Was Too Narrow: Mechanism Width as a Precondition for Interpretability-Driven Discovery** 2026  
(first author)

NeurIPS 2026 Workshop (Interp4Discovery), Atlanta

- The field’s standard intervention unit (one direction, one layer, one attention head) recovers only about **10%** of the achievable causal effect across three unrelated testbeds: rank-1 steering in vision encoders (0.111), single-layer ablation in Llama-3.1-8B (0.096), single-head sufficiency in GPT-2’s IOI circuit (0.086). Mechanism width  $k^*$  is **measurable before intervening**, and matching it recovers **95 to 100%**. Across 12 language models a cliff in the writer-score ladder predicts whether a single-unit test will work (Spearman  $\rho = 0.81$ ). (*with B. Joshi.*)
- Reframes negative interpretability results: a null at standard width is an **underpowered test, not evidence of absence**, and a cheap protocol (sweep width, compare  $k_{\text{repr}}^*$  to  $k_{\text{eff}}^*$ , run matched controls) separates the two before a discovery claim is made.

**Geometric Phases of Mechanism Formation in Neural Networks** 2026  
(sole-authored)

GLOW @ IJCAI-ECAI 2026 (Best Paper Candidate)

- Using linear probes and CKA across dense training checkpoints, shows classification mechanisms form **output-layer-first** and within the first  $\sim 5\%$  of training (Cohen’s  $d = 3.68$ ); the deep-first pattern holds in from-scratch LLM pretraining (GPT-2 Small, SmoLM2-135M) and reproduces on public Pythia and OLMo-2 checkpoints.

**MaiBERT: A Pre-training Corpus and Language Model for Low-Resourced Maithili** 2026  
LoResLM @ EACL 2026

- First monolingual BERT for Maithili ( $\sim 50\text{M}$  speakers); **87.02%** news-classification accuracy, outperforming multilingual baselines (MuRIL, NepBERTa), addressing fairness and access for underrepresented communities.

**How Far Can a Single Vector Carry a Language? Mechanistic Limits of Inference-Time Steering for Low-Resource Devanagari Languages** 2026  
(first author)

Under review, Machine Learning (Springer)

- With **no fine-tuning**, redirects a model’s **Hindi** representations toward unsupported Devanagari languages (Nepali, Maithili, Bhojpuri). Finds a shared **hourglass geometry** and an **adherence-fluency frontier** that bounds single-direction steering: adherence arrives before fluency, and **separability rather than proximity** predicts steerability (Hindi 3.23 vs English 1.37). Unifies the SafeConstellations and maiBERT lines.

**On the Relationship Between Representation Geometry and Generalization** 2026  
(sole-authored)

Pre-Print 2026

- Establishes that a label-free geometric metric, **effective dimension**, strongly predicts generalization across vision and language (partial  $r = 0.75$  over 52 ImageNet models, 13 architectures; carries to NLP encoders and decoder-only LLMs). The link is causal: degrading the geometry hurts accuracy while PCA cleanup preserves it, a diagnostic relevant to distribution-shift detection and robustness.

**The Steerability Spectrum of Visual Attributes** ongoing 2026

- Measures how readily different visual features (size, distance, orientation) can be steered inside vision models (DINOv2, CLIP, SigLIP): some lie on a single clean, controllable direction while others are distributed and resist steering, probing what accounts for the difference.

Also: **Revolutionizing Currency Security with YOLOv8** (J. Bus. Econ. Stud., 2024; 0.986 true-positive recall) · **Machine Learning Analysis of Tirhuta Lipi** (2023) · **Support Vectors for Imbalanced Text Classification** (2023).

## EXPERIENCE

**Astha.ai**  
AI Researcher, AI Security & Agentic Systems

USA (Remote)  
May 2025 – Present

- Research-to-production work in AI security: **primary author** of the SAFE-MCP vulnerability scanner (**78+** MITRE ATT&CK-mapped techniques covering tool poisoning and injected tool descriptions) and **lead engineer** on a Go-based Kubernetes compliance agent running both inside customer clusters and in the SaaS platform.
- Recurring research thread, **making LLM-based verdicts reproducible**: a statistical analysis mode for the MCP scanner and a repeated-run consensus mechanism with fixed sampling parameters for compliance verdicts, alongside LLM-driven scan planning over a catalog of **~470** compliance-control skills and prompt-injection hardening with tree-sitter semantic chunking.

#### AMNIL Technologies

AI Engineer, RAG & Infrastructure

Nepal (Hybrid)

Jun 2024 – May 2025

- Guardrails + **LLM-as-a-Judge** evaluation; hybrid-search RAG (Qdrant, +35% retrieval accuracy); self-hosted LLM serving with vLLM (40% lower latency).

#### GradeUp Educations

Data Team Lead

Kathmandu, Nepal

Jan 2022 – Jun 2024

- Led the data team: learning agents/chatbots, an automated grade-evaluation system, and semantic-similarity matching (NLP + image processing).

#### DeepLearning.AI

GAN Specialization Mentor

Global (Remote)

Aug 2021 – Present

- Technical mentor for the GANs Specialization (loss functions, training stability, generative-model robustness).

### SELECTED PROJECTS & OPEN SOURCE

**Removing the “Workspace” Layers of Llama**: interpretability study locating a decodable-yet-entangled middle band of layers in Llama-3.1-8B; ablating the whole band collapses reasoning while shallow recall survives (linear probes + reversible layer-ablation, nnsight/NDIF to 405B). [code].

**World Models (toy experiments)**: built Dreamer-style world models (RSSM: GRU + stochastic latents) from scratch for continuous and visual control; agents trained purely in imagination transfer to the real task (CartPole 500/500), and short-horizon dreams with periodic reality-grounding cut drift 11×. [code].

**SAFE-MCP**: community security framework for the Model Context Protocol (ATT&CK-style threat/defense rules; prompt-injection & tool-poisoning modules).

**Telebot-Claude-Bridge**: permission-gated, human-in-the-loop remote control of the Claude Code CLI (inline allow/deny per tool call, audit logging).

**Maithili speech**: open Whisper-tiny **ASR** and MMS-TTS **TTS** for Maithili (live HF Spaces demos); **maiBERT**, the first open BERT for Maithili (HF).

**NPL Coders**: founder of Nepal’s national data-science competition community.

### TECHNICAL PROFICIENCY

**Interpretability & Alignment**: representation & activation steering, mechanistic interpretability, linear probes / CKA, representation geometry, effective dimension, over-refusal mitigation.

**LLMs & NLP**: LLMs (Llama, Claude, GPT), fine-tuning (PEFT/LoRA), RAG, multilingual / low-resource NLP, tokenizer training, embedding alignment.

**Agentic & Security**: Model Context Protocol, multi-agent frameworks, MITRE ATT&CK mapping, prompt-injection & tool-poisoning defense, adversarial evaluation.

**Tools**: Python (expert), PyTorch, TensorFlow, vLLM, FastAPI, Docker, Kubernetes, vector DBs (Qdrant, ChromaDB), nnsight.

**Research methods**: empirical NLP evaluation, information-theoretic metrics, adversarial benchmarking, statistical analysis.

### EDUCATION

#### Pulchowk Engineering College, IOE, Tribhuvan University

Bachelor of Computer Engineering

Kathmandu, Nepal

2019 – 2024 (conferred 2026)

- B.E. thesis: “Evaluating Auto-Encoding Transformer Language Models for Maithili Text Classification”: my first Maithili masked language model, and the foundation for maiBERT.
- Key subjects: Artificial Intelligence, Data Structures & Algorithms, Distributed Systems, Big Data Technologies, Network Security, Compiler Design.

### HONORS & AWARDS

**Best Paper Candidate, GLOW @ IJCAI-ECAI 2026** (sole-authored) · Outstanding Reviewer, Mechanistic Interpretability Workshop (ICML 2026) · Winner, GritFeat AI Hackathon (2023) · 1st Runner-Up, Locus Dataverse (2023) · Best AI Project, DELTA 3.0 (2023) · Winner, IT-Meet Image Challenge (2022) · Winner, LogPoint Capture The Flag (2022)